Document No: BGen/2025-10/HR/Technical/Recruitad/001

Generative AI – Engineer

Job Summary:

We are looking for a Generative AI - Engineer to join our growing team to help design, build, fine-tune, and deploy cutting - edge generative AI models and agentic systems. You will work on the full lifecycle of foundational model development - involving both large and small language models (LLMs and SLMs) - to create scalable AI solutions that address diverse needs across different business domains. This role is ideal for proactive individuals with a strong foundation in machine learning and an experimental mindset, who are passionate about driving transformative advancements in generative AI from research to real-world production impact.

Key Responsibilities:

- Develop and train foundational generative AI models across modalities such as text-to-text, text-to-speech, automatic speech recognition, and vision language.
- Fine-tune and adapt models for specific tasks and domains.
- Build and maintain pipelines for data curation, preprocessing, training, evaluation, and continuous improvement of models.
- Implement debugging, CI/CD, and observability to ensure reliability and efficiency across the development lifecycle.
- Develop retrieval-augmented generation (RAG) pipelines and optimize prompt engineering strategies.
- Optimize training and inference performance through quantization, distributed training/inference, GPU/TPU acceleration.
- Monitor, benchmark, and improve model performance with a focus on accuracy, efficiency, and reducing hallucinations.
- Collaborate with cross-functional teams to build robust AI stacks and integrate them seamlessly into production pipelines for deployment.
- Document technical processes, AI model architectures, and experimental results, while maintaining well-structured, version-controlled code repositories.
- Stay current with advancements in transformer architectures, open-source releases, and AI tooling.

Minimum Qualifications and Experience:

- Bachelor's or Master's in Computer Science, AI/ML, Data Science or any related field with 2 to 5 years of industry experience in applied machine learning or AI development.

Required Expertise:

- Proficiency in Python programming with solid foundation in computer science fundamentals such as data structures and algorithms.
- Strong problem-solving skills and demonstrated ability to lead projects.
- Hands-on experience with a few of the tools listed below:

- One or more model libraries and ML frameworks such as TensorFlow, PyTorch, HF Transformers, NeMo, etc.
- AI application libraries and orchestration frameworks such as DSPy, Langgraph, Langchain, Llamaindex, etc.
- GPU/TPU based training and inference using libraries such as vLLM.
- Distributed training tools such as SLURM, Ray, Pytorch DDP, NCCL, etc.
- Version control, observability systems, and MLOps tools such as Git, DVC, W&B, MLFlow, KubeFlow, etc.
- Data analysis and curation tools such as Dask, Milvus, Apache Spark, Numpy, etc.
- Chunking, embeddings, vector databases (e.g., Pinecone, Weaviate, Milvus), and retrieval-augmented generation (RAG).
- Model context protocol (MCP), Agent to Agent (A2A), and Agent Communication Protocol (ACP).

- Team player with excellent interpersonal skills and ability to collaborate effectively with remote team members.
- Go-getter attitude and ability to flourish in a fast-paced, startup environment.
- Prior experience of building and deploying LLMs or SLMs, experience with multimodal models, and track record of contributions to open-source AI/ML projects would be a big plus.